

Reinforcement Learning with Deep Q-Networks for Classical Arcade Games

Gabriel Ewing
Department of Electrical Engineering
and Computer Science
Case Western Reserve University
Cleveland, OH
gre5@case.edu

Abstract—I implemented deep neural networks for playing Atari 2600 games using Q-Learning. We create Voronoi diagrams for two-dimensional maps with point obstacles. We explore the algebraic and geometric intricacies of the Sweepline algorithm used to build the diagrams. We review the concept of a k-Nearest-Neighbors classifier and use the completed Voronoi diagram to find nearest neighbors.

I. INTRODUCTION

Over the past few decades, computers have been able to play a series of increasingly complex games at a high level, perhaps most notably in Deep Blue’s defeat of reigning world chess champion Garry Kasparov in 1997 and AlphaGo’s victory against top Go player Lee Sedol in 2016. Complexity, in this case, is mostly tied to the *branching factor* of the game in question; that is, the number of actions that are available to a player at a point in time. While chess and Go are both complex beyond the point of human understanding, chess was won much sooner by computers because it has a significantly lower branching factor than Go.

Despite their massive possible game trees, both of the aforementioned games are laughably simple compared to the real world in which we live every day. This is in part due to the fact that the state at any given time for both games can be entirely encoded by a relatively small number of bits, whereas the real world has meaningful state down to the atomic level and a functionally infinite branching factor. The low-dimensional game representation means that vision, one of the most elaborate functions of our brains, has little to no bearing on our ability to play either game.

The natural next targets for computational game-playing, then, are games that require some visual processing to play effectively. In 2013, Bellemare et al. published the Arcade Learning Environment (ALE) [1], which allows programmatic access to an emulator for video games on the Atari 2600 console.

Note that in this paper we mainly follow Mnih et al. [2], and omit some details from that paper for brevity, but attempt to expand on areas that were given less attention in that paper and we found interesting during our implementation.

A. Q-Learning

Parallel to the trend of increasing complexity, focus has shifted from explicitly teaching computers to play games to having the computers learn for themselves from experience. This is in part due to stronger general interest in machine learning, and in part due to the learning approach surpassing the teaching approach.

The subfield of machine learning most relevant to gameplay is called reinforcement learning. Reinforcement learning is the problem of solving Markov Decision Processes (MDPs). An MDP is a task consisting of a set of states, a set of actions, a set of transition probabilities from the set of (state, action) pairs to the set of outcome states, a set of starting states, and a reward function. At each step of the MDP, the agent observes a state, selects an action, and is given a reward signal. The reward signal is what the agent should try to optimize.

A prominent classical reinforcement learning approach is Q-Learning. In Q-Learning, the agent maintains a function that represents the value of a given (state, action) pair: that is, the expected cumulative reward received by taking an action in a state. If the agent has learned a good Q-function, it can then exploit that by taking the action with the maximal Q-value in a given state.

Usually, the Q-function is parameterized by a linear combination of the state features. Its weights are updated by the *Bellman equation*:

$$Q_{new}(s, a) \leftarrow \alpha \left(R(s, a) + \gamma \max_{a'} Q(s', a') \right) \quad (1)$$

α is the *learning rate*, which controls how much the weights move at each update. $R(s, a)$ is the reward received after taking action a in state s . γ is the *discount factor*, representing how much less a reward at a future step is worth than a reward at the current step. s' is the state transitioned to after taking action a .

Rather than use a set of linear weights for the Q-function, Deep Q-Learning uses a represents the Q-function via a neural network. We will explore this concept further in Section II-A.

B. Arcade Learning Environment

The ALE provides an interface to game binaries for Python and Java (binaries must be acquired from third-party sites).



Fig. 1: The Breakout Atari game. The goal is to bounce the ball with the bar and hit all the tiles, while avoiding letting the ball drop below the screen.

After a game is initialized, play proceeds in a loop of the agent sending an action, receiving a reward, and requesting the frame of the current game state.

Frames are given as 210×160 pixel images with 128 colors. To reduce the computational load, Mnih et al. [2] downsampled the frame to 110×84 and converted the colors to grayscale. We also adopted this approach. Conveniently, the ALE has an API call to get a grayscale frame instead of RGB.

The actions available to an agent generally include up, down, left, right, and game-specific actions such as jump or shoot. These are encoded as integers by the ALE, so the agent has no semantic knowledge of the actions before it starts playing the game.

II. METHODS

A. Deep Q-Learning

As mentioned in Section I-A, Deep Q-Learning uses a neural network to represent the Q-function. This formulation has a couple benefits. First, deep neural networks are capable of learning extremely complex value functions, whereas linearly-parameterized functions are limited by the complexity of the input features. This is important for any game with a large state space, such as those we are exploring here. Second, deep neural networks have been the driving force of the recent revolutionary improvements in computer vision. For games with visually-represented states, we can take advantage of those improvements.

There are some differences between the use and updating of a deep Q-network and a linearly-parameterized Q-function.

1) *Action Representation*: One possible network architecture would be to have a separate Q-network for each possible action. We could then feed a state through each network and select the action with the highest Q-value. This approach would work, but it ignores some possible performance improvements and recent discoveries in neural networks.

Neural networks are computationally expensive to train, and are even more expensive when the number of weights in the network is high. This means that when possible, we should try to construct our architecture so that the number of weights that need to be updated is minimized. One way to do this is to only use one network for all of the actions, and have each action correspond to one of the network’s output nodes.

Besides the computational improvements, this takes advantage of the fact that the early-layer processing for each action should be virtually identical, since they all require knowledge of the relevant portions of the game state derived from the visual input. By using shared early layers, we ensure that useful visual processing patterns that would have only been “discovered” by one network in a separate-action architecture are instead shared by all of the actions.

This is similar to the principal that we saw in class where early-layer weights for image classification tasks can usually be reused effectively, even for classification on images that would seem unrelated to those that produced the original weights.

2) *Network Feedback*: The Bellman equation, discussed in Section I-A, is a *prescriptive* update equation; that is, it gives the steps that should be taken to update the Q-function. When we have a linear set of weights, the Bellman equation is sufficient and complete because we know exactly how the outputs will change with any alteration of the weights.

In contrast, with a Q-network, we need a *descriptive* update formulation. Since the backpropagation equations for neural networks are established and widely available, we need to *tell the network what it should produce* and then train it via backpropagation to match that output.

Mnih et al. [2] found a way to derive an error value for backpropagation given a reward and a transition:

$$Err(r, s, a, s') = \left(r + \gamma \max_{a'} Q(s', a') - Q(s, a) \right)^2 \quad (2)$$

This still uses the Bellman contraction from a reward and future Q-value to figure out how “wrong” the current Q-value is.

B. Network Architecture

We mostly follow the architecture presented in [2], with some minor alterations.

1) *Downsampling*: As mentioned previously, they downsampled the input frames to 110×84 . Furthermore, due to their reliance on a specific GPU convolution implementation that required square inputs, they truncated the sides of each downsampled frame so that the resulting output was 84×84 . With modern neural network libraries, this is unnecessary, so we avoid the truncation. However, to make the modular arithmetic work nicely, we downsample to 108×84 rather than 110×84 .

2) *Frame Stacking*: To reduce the number of network operations per game step, they stacked four downsampled frames simultaneously and used that as the input to the network. This means that the network only is called to get

a new action every four game steps, and the agent simply repeats the previous action in between. We adopt this strategy.

3) *Convolutional Layers:* There are two convolutional layers in the Q-network: first, a layer of 16 8×8 convolution patches with stride 4, followed by a layer of 32 4×4 convolution patches with stride 2. Each uses TensorFlow's rectified linear unit activation function (`tf.nn.relu`).

4) *Fully-Connected:* Following the convolutional layers is a fully-connected layer with 128 units, also using the rectified linear activation function. The fully-connected layer is connected to the output layer, which has a number of nodes corresponding to the number of actions available in the given game.

C. Training

1) *Experience Replay:* Classical Q-Learning updates the Q-function with the last transition seen by the MDP. This would probably work in the deep Q-Learning paradigm, but the authors in [2] found what is likely an improvement. Instead of throwing away the transitions after an update is performed, they store a large number of the most recent transitions. At each update step, they sample from this memory and update the network using the sampled transitions. The main benefit of this is that the network does not deviate too much when the state goes in one direction for an extended period of time.

We discovered that each recorded transition requires quite a bit of data: the prior frame stack, the next frame stack, the action taken, and the average reward. Just storing the Q-values is insufficient, because the network update is supposed to be done with the error from the current weights, not a previous set of weights. Additionally, we have to store the action taken, because the action that maximizes the Q-value may have changed since the transition was recorded.

2) *Backpropagation:* We used the RMSProp algorithm, similarly to [2].

III. RESULTS

Unfortunately, we were not able to generate meaningful results due to performance issues. In [2], they trained each game for 10,000,000 steps. With frame stacking, that means there were 2,500,000 network updates. When we tried to run our implementation with TensorFlow, it took hours to reach 100 steps even on a powerful computer. This confused us for quite a while: surely the previous authors did not have access to hardware five orders of magnitude faster than ours.

The problem was that they were not training the network to convergence on the input transitions at every step. This was not discussed in their paper, but it seems like the only reasonable explanation.

Unfortunately, TensorFlow does not have a way to terminate one of their out-of-the-box optimization algorithms early without stopping the entire program. This means that a TensorFlow implementation of this algorithm that has a plausible chance at working in a reasonable amount of time needs to be implemented using their direct gradient computation API, which is significantly more involved than calling

one of the optimization algorithms. We discovered this too late to make a meaningful amount of progress on the alternate implementation.

There is a small plus to training to convergence, which is that it does more improvement to the network per step. However, the problem we ran into was that there were so few game steps taken that the agent did not have time to learn anything about the game or do sufficient exploration.

IV. CONCLUSION

Not getting any results was obviously disappointing. Additionally, we didn't get very far beyond implementing the 2013 paper. We did learn quite a bit about using TensorFlow and deep Q-Learning, and hopefully provided some interesting insights from the challenges we encountered.

REFERENCES

- [1] M. G. Bellemare, Y. Naddaf, J. Veness, and M. Bowling, "The arcade learning environment: An evaluation platform for general agents," *Journal of Artificial Intelligence Research*, vol. 47, pp. 253–279, 06 2013.
- [2] V. Mnih, K. Kavukcuoglu, D. Silver, A. Graves, I. Antonoglou, D. Wierstra, and M. Riedmiller, "Playing atari with deep reinforcement learning," *arXiv preprint arXiv:1312.5602*, 2013.
- [3] V. Mnih, K. Kavukcuoglu, D. Silver, A. A. Rusu, J. Veness, M. G. Bellemare, A. Graves, M. Riedmiller, A. K. Fidjeland, G. Ostrovski *et al.*, "Human-level control through deep reinforcement learning," *Nature*, vol. 518, no. 7540, pp. 529–533, 2015.